

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score
import warnings
warnings.filterwarnings('ignore')
```

```
In [4]: df = pd.read_csv(r'C:\Users\HP\Downloads\IRIS.csv')
```

```
In [5]: df
```

```
Out[5]:
```

| | sepal_length | sepal_width | petal_length | petal_width | species |
|------------|--------------|-------------|--------------|-------------|----------------|
| 0 | 5.1 | 3.5 | 1.4 | 0.2 | Iris-setosa |
| 1 | 4.9 | 3.0 | 1.4 | 0.2 | Iris-setosa |
| 2 | 4.7 | 3.2 | 1.3 | 0.2 | Iris-setosa |
| 3 | 4.6 | 3.1 | 1.5 | 0.2 | Iris-setosa |
| 4 | 5.0 | 3.6 | 1.4 | 0.2 | Iris-setosa |
| ... | ... | ... | ... | ... | ... |
| 145 | 6.7 | 3.0 | 5.2 | 2.3 | Iris-virginica |
| 146 | 6.3 | 2.5 | 5.0 | 1.9 | Iris-virginica |
| 147 | 6.5 | 3.0 | 5.2 | 2.0 | Iris-virginica |
| 148 | 6.2 | 3.4 | 5.4 | 2.3 | Iris-virginica |
| 149 | 5.9 | 3.0 | 5.1 | 1.8 | Iris-virginica |

150 rows × 5 columns

```
In [6]: df.isna().sum()
```

```
Out[6]: sepal_length    0
sepal_width          0
petal_length         0
petal_width          0
species              0
dtype: int64
```

```
In [7]: df.duplicated().sum()
```

```
Out[7]: 3
```

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   sepal_length    150 non-null    float64
1   sepal_width     150 non-null    float64
2   petal_length    150 non-null    float64
3   petal_width     150 non-null    float64
4   species         150 non-null    object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
In [9]: df.describe()
```

```
Out[9]:
```

| | sepal_length | sepal_width | petal_length | petal_width |
|-------|--------------|-------------|--------------|-------------|
| count | 150.000000 | 150.000000 | 150.000000 | 150.000000 |
| mean | 5.843333 | 3.054000 | 3.758667 | 1.198667 |
| std | 0.828066 | 0.433594 | 1.764420 | 0.763161 |
| min | 4.300000 | 2.000000 | 1.000000 | 0.100000 |
| 25% | 5.100000 | 2.800000 | 1.600000 | 0.300000 |
| 50% | 5.800000 | 3.000000 | 4.350000 | 1.300000 |
| 75% | 6.400000 | 3.300000 | 5.100000 | 1.800000 |
| max | 7.900000 | 4.400000 | 6.900000 | 2.500000 |

```
In [10]: df.drop_duplicates(inplace=True)
```

```
In [11]: df
```

```
Out[11]:
```

| | sepal_length | sepal_width | petal_length | petal_width | species |
|------------|--------------|-------------|--------------|-------------|----------------|
| 0 | 5.1 | 3.5 | 1.4 | 0.2 | Iris-setosa |
| 1 | 4.9 | 3.0 | 1.4 | 0.2 | Iris-setosa |
| 2 | 4.7 | 3.2 | 1.3 | 0.2 | Iris-setosa |
| 3 | 4.6 | 3.1 | 1.5 | 0.2 | Iris-setosa |
| 4 | 5.0 | 3.6 | 1.4 | 0.2 | Iris-setosa |
| ... | ... | ... | ... | ... | ... |
| 145 | 6.7 | 3.0 | 5.2 | 2.3 | Iris-virginica |
| 146 | 6.3 | 2.5 | 5.0 | 1.9 | Iris-virginica |
| 147 | 6.5 | 3.0 | 5.2 | 2.0 | Iris-virginica |
| 148 | 6.2 | 3.4 | 5.4 | 2.3 | Iris-virginica |
| 149 | 5.9 | 3.0 | 5.1 | 1.8 | Iris-virginica |

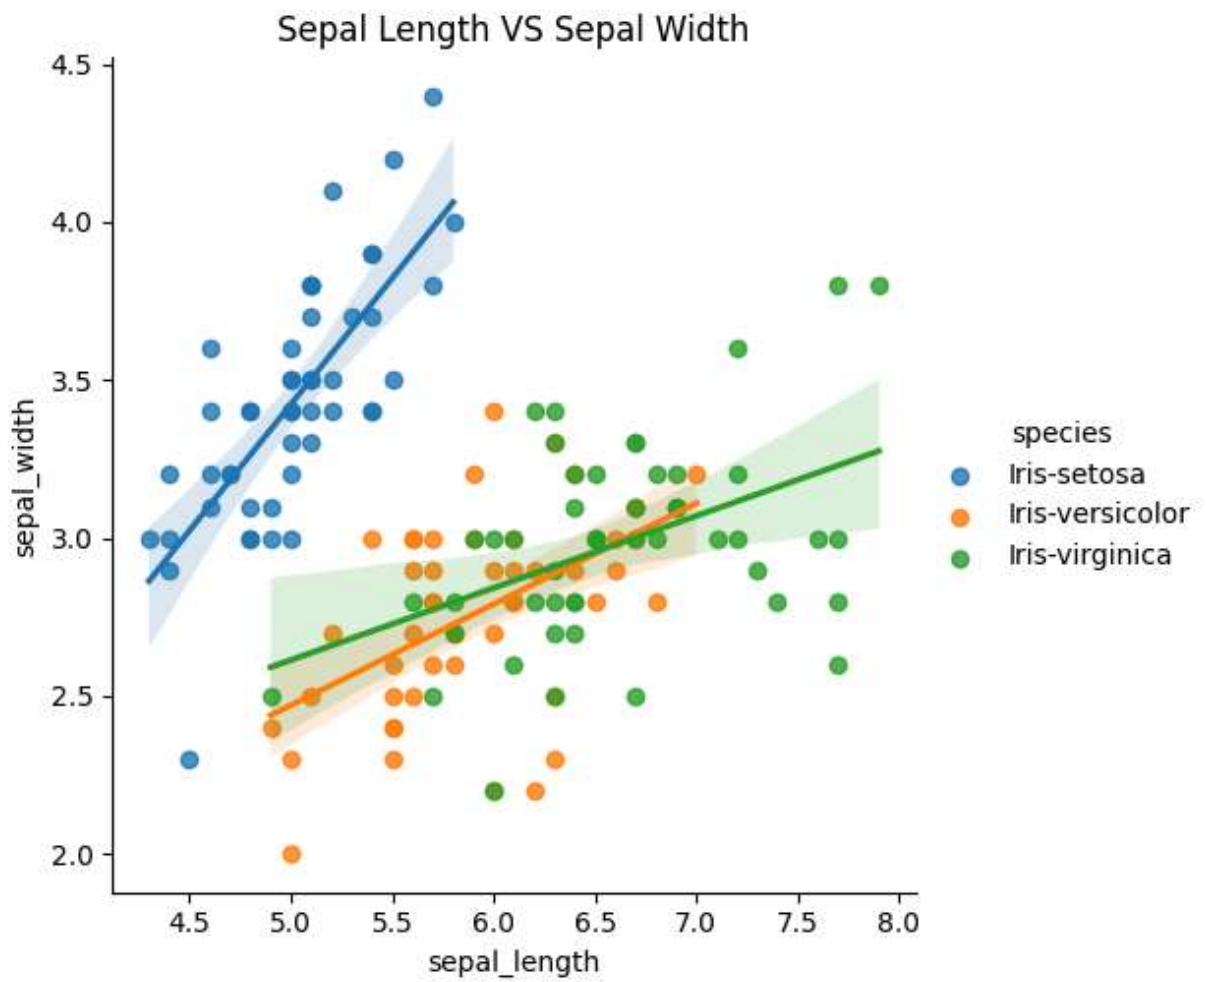
147 rows × 5 columns

```
In [12]: df['species'].value_counts()
```

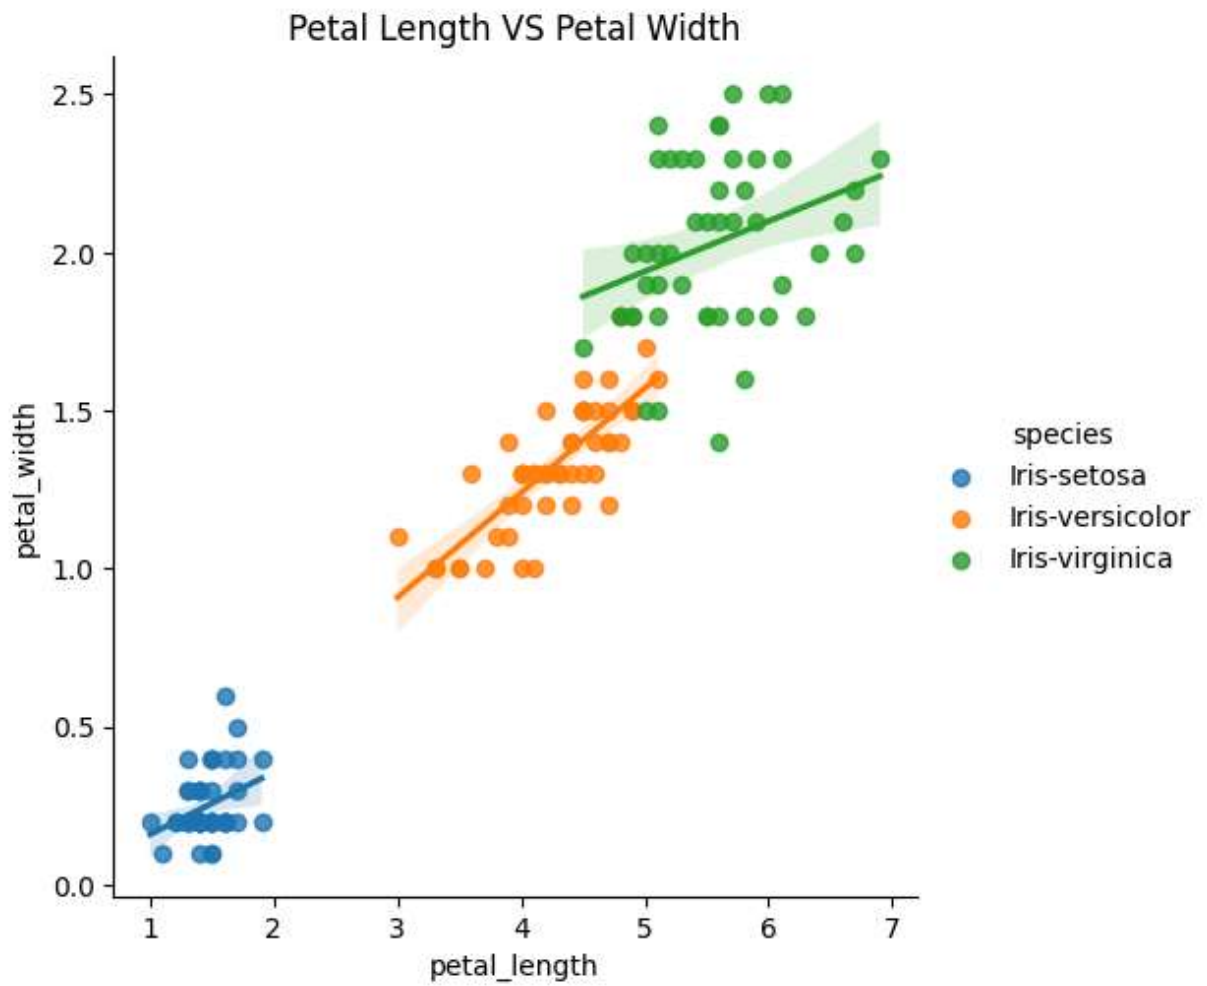
```
Out[12]: species
Iris-versicolor    50
Iris-virginica     49
Iris-setosa        48
Name: count, dtype: int64
```

```
In [13]: sns.lmplot(
    x="sepal_length",
    y="sepal_width",
    hue="species",
    data=df
)

plt.title("Sepal Length VS Sepal Width")
plt.show()
```



```
In [14]: sns.lmplot(  
    x="petal_length",  
    y="petal_width",  
    hue="species",  
    data=df  
)  
  
plt.title("Petal Length VS Petal Width")  
plt.show()
```



```
In [15]: label_encoder = LabelEncoder()  
df['species'] = label_encoder.fit_transform(df['species'])
```

```
In [16]: df
```

```
Out[16]:
```

| | sepal_length | sepal_width | petal_length | petal_width | species |
|------------|--------------|-------------|--------------|-------------|---------|
| 0 | 5.1 | 3.5 | 1.4 | 0.2 | 0 |
| 1 | 4.9 | 3.0 | 1.4 | 0.2 | 0 |
| 2 | 4.7 | 3.2 | 1.3 | 0.2 | 0 |
| 3 | 4.6 | 3.1 | 1.5 | 0.2 | 0 |
| 4 | 5.0 | 3.6 | 1.4 | 0.2 | 0 |
| ... | ... | ... | ... | ... | ... |
| 145 | 6.7 | 3.0 | 5.2 | 2.3 | 2 |
| 146 | 6.3 | 2.5 | 5.0 | 1.9 | 2 |
| 147 | 6.5 | 3.0 | 5.2 | 2.0 | 2 |
| 148 | 6.2 | 3.4 | 5.4 | 2.3 | 2 |
| 149 | 5.9 | 3.0 | 5.1 | 1.8 | 2 |

147 rows × 5 columns

```
In [19]: x = df.drop(columns='species')
         y = df.species
```

```
In [18]: x
```

```
Out[18]:
```

| | sepal_length | sepal_width | petal_length | petal_width |
|------------|--------------|-------------|--------------|-------------|
| 0 | 5.1 | 3.5 | 1.4 | 0.2 |
| 1 | 4.9 | 3.0 | 1.4 | 0.2 |
| 2 | 4.7 | 3.2 | 1.3 | 0.2 |
| 3 | 4.6 | 3.1 | 1.5 | 0.2 |
| 4 | 5.0 | 3.6 | 1.4 | 0.2 |
| ... | ... | ... | ... | ... |
| 145 | 6.7 | 3.0 | 5.2 | 2.3 |
| 146 | 6.3 | 2.5 | 5.0 | 1.9 |
| 147 | 6.5 | 3.0 | 5.2 | 2.0 |
| 148 | 6.2 | 3.4 | 5.4 | 2.3 |
| 149 | 5.9 | 3.0 | 5.1 | 1.8 |

147 rows × 4 columns

```
In [20]: y
```

```
Out[20]: 0      0
          1      0
          2      0
          3      0
          4      0
          ..
        145     2
        146     2
        147     2
        148     2
        149     2
        Name: species, Length: 147, dtype: int32
```

```
In [21]: X_train, x_test, Y_train, y_test = train_test_split(x, y, test_size=0.2, random_sta
```

```
In [24]: from sklearn.linear_model import LogisticRegression
```

```
# Instantiate the model
model = LogisticRegression()

# Fit the model with training data
model.fit(X_train, Y_train)
```

```
Out[24]: LogisticRegression ⓘ ?
         LogisticRegression()
```

```
In [25]: model.score(x_test, y_test)
```

```
Out[25]: 1.0
```

```
In [26]: model.score(X_train, Y_train)
```

```
Out[26]: 0.9743589743589743
```

```
In [27]: model.predict([[5.1,3.5,1.4,0.2]])
```

```
Out[27]: array([0])
```

```
In [28]: y_predicted = model.predict(x_test)
```

```
In [29]: y_predicted
```


```
Out[29]: array([2, 1, 1, 1, 0, 0, 1, 0, 2, 2, 0, 2, 0, 2, 1, 2, 2, 0, 1, 1, 0, 2,
                1, 0, 2, 0, 0, 2, 2, 0])
```

```
In [30]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_predicted) # compare between predicted values, actu
cm
```

```
Out[30]: array([[11,  0,  0],
                [ 0,  8,  0],
                [ 0,  0, 11]], dtype=int64)
```

```
In [31]: import seaborn as sn
plt.figure(figsize = (10,7))
sn.heatmap(cm, annot = True)
plt.xlabel('Predicted')
plt.ylabel('Truth')
```

Out[31]: Text(95.7222222222221, 0.5, 'Truth')



In []: